

```
In [180]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.preprocessing import OrdinalEncoder
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import StandardScaler

from sklearn.model_selection import train_test_split
from sklearn.model_selection import validation_curve
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV

from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso

from sklearn.svm import SVR
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import AdaBoostRegressor

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.svm import SVC

from sklearn.decomposition import PCA
from sklearn.cluster import KMeans

In [2]: # Load Data
df = pd.read_csv("5241dataset.csv", encoding = 'unicode_escape')

/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/1370
471252.py:2: DtypeWarning: Columns (27,41) have mixed types. Specify
dtype option on import or set low_memory=False.
df = pd.read_csv("5241dataset.csv", encoding = 'unicode_escape')
```

1. Dataset Description

```
In [3]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 34931 entries, 0 to 34930
Data columns (total 49 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   Record number                        34931 non-null  int64

1   Individual ID                        33200 non-null  float64
2   Predator                            34931 non-null  object
3   Predator common name                 34931 non-null  object
4   Predator taxon                       34931 non-null  object
5   Predator lifestage                   34931 non-null  object
6   Type of feeding interaction           34931 non-null  object
7   Predator length                      34931 non-null  float64
8   Predator length unit                  34931 non-null  object
9   Predator standard length              34930 non-null  float64
10  Predator total length                 34818 non-null  float64
11  Predator TL/FL/SL conversion reference 29683 non-null  object
12  Standardised predator length          34931 non-null  float64
13  Predator mass                         34931 non-null  float64
14  Predator mass unit                    34931 non-null  object
15  Predator mass check                   34931 non-null  float64
16  Predator mass check diff               34931 non-null  float64
17  Predator ratio mass/mass              34931 non-null  float64
18  SI predator mass                      34931 non-null  float64
19  Diet coverage                         34931 non-null  object
20  Prey                                  34931 non-null  object
21  Prey common name                     34931 non-null  object
22  Prey taxon                           34931 non-null  object
23  Prey length                          34931 non-null  float64
24  Prey length unit                      34929 non-null  object
25  SI prey length                       34931 non-null  float64
26  Prey width                           964 non-null    float64
27  Prey width unit                       964 non-null    object
28  Prey mass                            34931 non-null  float64
29  Prey mass unit                       34931 non-null  object
30  Prey mass check                      34931 non-null  float64
31  Prey mass check diff                  34931 non-null  float64
32  Prey ratio mass/mass                  34931 non-null  float64
33  SI prey mass                          34931 non-null  float64
34  Geographic location                  34931 non-null  object
35  Latitude                             34931 non-null  object
36  Latitude Degree                       34931 non-null  int64
37  Latitude Minute                       34931 non-null  int64
38  Latitude label                        34931 non-null  object
39  Longitude                             34931 non-null  object
40  Longitude Degree                      34931 non-null  int64
41  Longitude Minute                      34931 non-null  object
42  Longitude label                       34931 non-null  object
43  Depth                                34931 non-null  int64
44  Mean annual temp                      34931 non-null  float64
45  SD annual temp                        34931 non-null  float64
46  Mean PP                              34931 non-null  int64
47  SD PP                                34931 non-null  float64
48  Specific habitat                      34931 non-null  object
dtypes: float64(21), int64(6), object(22)
memory usage: 13.1+ MB

In [4]: df.describe().T

Out[4]:
```

	count	mean	std	min	25%	50%	
Record number	34931.0	17466.000000	10083.855463	1.000000e+00	8733.50000	17466.0000	26
Individual ID	33200.0	7767.209699	5603.773065	1.000000e+00	2066.75000	7000.0000	12
Predator length	34931.0	411.708758	865.149724	3.000000e+00	67.50000	200.0000	
Predator standard length	34930.0	275.498247	851.527130	0.000000e+00	18.50000	92.0000	
Predator total length	34818.0	410.125251	891.934132	0.000000e+00	73.41000	210.0000	
Standardised predator length	34931.0	68.362248	53.988482	3.000000e-01	28.92000	68.0000	
Predator mass	34931.0	15292.363560	46260.898001	1.140000e-04	172.01000	1751.1000	6
Predator mass check	34931.0	57433.601425	142121.384771	1.413000e-03	1265.80000	16455.0000	43
Predator mass check diff	34931.0	42141.261535	97206.742764	1.298700e-03	995.64000	12997.0000	37
Predator ratio mass/mass	34931.0	7.488458	5.037282	1.053500e+00	4.11160	6.5149	
SI predator mass	34931.0	15292.363560	46260.898001	1.140000e-04	172.01000	1751.1000	6
Prey length	34931.0	72.617237	214.978864	8.000000e-02	4.68225	21.9890	
SI prey length	34931.0	8.128151	9.084534	4.160000e-04	2.19945	5.0000	
Prey width	964.0	266.169233	238.438741	4.530000e-01	110.00000	216.0000	
Prey mass	34931.0	28.409896	121.990812	7.530000e-11	0.16240	1.4888	
Prey mass check	34931.0	264.665719	2501.661458	3.770000e-12	0.55684	6.5417	
Prey mass check diff	34931.0	236.258016	2460.368288	-1.151000e+02	0.22081	4.6817	
Prey ratio mass/mass	34931.0	8.610682	35.534542	1.915600e-02	3.19750	5.2333	
SI prey mass	34931.0	28.407359	121.991367	7.530000e-11	0.16240	1.4888	
Latitude Degree	34931.0	40.118605	12.041278	8.000000e+00	40.00000	40.0000	
Latitude Minute	34931.0	9.778821	21.456202	0.000000e+00	0.00000	0.0000	
Longitude Degree	34931.0	47.276545	28.059336	2.000000e+00	14.00000	70.0000	
Depth	34931.0	1561.806418	1768.576051	5.000000e+00	677.00000	677.0000	3
Mean annual temp	34931.0	15.166128	6.466898	-1.300000e+00	13.90000	15.1000	
SD annual temp	34931.0	4.366302	2.285000	6.000000e-01	2.50000	5.6000	
Mean PP	34931.0	727.152243	320.066435	9.100000e+01	437.00000	867.0000	
SD PP	34931.0	69.225888	34.382417	7.000000e+00	38.00000	80.0000	

2. Data Cleaning

2.1 Missing Values

```
In [5]: col_na = df.columns[df.isna().any()]
col_na_ratio = df.isna().sum()/df.shape[0]


print("The columns contain missing values are:\n", col_na)

fig,ax = plt.subplots(figsize=(15,4))
sns.barplot(x = df.columns, y = col_na_ratio, ax = ax)

ax.tick_params(axis = 'x', rotation = 90)
ax.set_xlabel('Features')
ax.set_ylabel('Missing Rate')
ax.set_title('Missing Rates of Different Features')

plt.tick_params(labelsize=6)
plt.show()
```

The columns contain missing values are:
Index(['Individual ID', 'Predator standard length', 'Predator total length',
'Predator TL/FL/SL conversion reference', 'Prey length unit',
'Prey width', 'Prey width unit'],
dtype='object')



```
In [6]: df.drop(['Individual ID','Predator TL/FL/SL conversion reference','Prey length unit',
```

```
In [7]: df.columns
```

```
Out[7]: Index(['Record number', 'Predator', 'Predator common name', 'Predator taxon',  

'Predator lifestage', 'Type of feeding interaction', 'Predator length',  

'Predator length unit', 'Predator standard length',  

'Predator total length', 'Standardised predator length',  

'Predator mass', 'Predator mass unit', 'Predator mass check',  

'Predator mass check diff', 'Predator ratio mass/mass',  

'SI predator mass', 'Diet coverage', 'Prey', 'Prey common name',  

'Prey taxon', 'Prey length', 'Prey length unit', 'SI prey length',  

'Prey mass', 'Prey mass unit', 'Prey mass check',  

'Prey mass check diff', 'Prey ratio mass/mass', 'SI prey mass',  

'Geographic location', 'Latitude', 'Latitude Degree', 'Latitude label',  

'Longitude', 'Longitude Degree', 'Longitude label', 'Longitude Minute',  

'Longitude label', 'Depth', 'Mean annual temp', 'SD annual temp',  

'Mean PP', 'SD PP', 'Specific habitat'],  

dtype='object')
```

2.2 Dataset Split

```
In [8]: df_predator = df[['Record number', 'Predator', 'Predator common name',  

'Predator lifestage', 'Type of feeding interaction', 'Predator length',  

'Predator length unit', 'Predator standard length', 'Predator total length',  

'Standardised predator length', 'Predator mass', 'Predator mass unit',  

'Predator mass check', 'Predator mass check diff', 'Predator ratio mass/mass',  

'SI predator mass', 'Diet coverage', 'Prey', 'Prey common name', 'Prey taxon',  

'Prey length', 'Prey length unit', 'SI prey length', 'Prey mass', 'Prey mass unit',  

'Prey mass check', 'Prey mass check diff', 'Prey ratio mass/mass', 'SI prey mass',  

'Geographic location', 'Latitude', 'Latitude Degree', 'Latitude label',  

'Longitude', 'Longitude Degree', 'Longitude label', 'Longitude Minute', 'Longitude label',  

'Depth', 'Mean annual temp', 'SD annual temp', 'Mean PP', 'SD PP', 'Specific habitat']]

df_prey = df[['Prey', 'Prey common name', 'Prey taxon', 'Prey length', 'Prey length unit',  

'Prey mass', 'Prey mass unit', 'Prey mass check', 'Prey mass check diff',  

'Prey ratio mass/mass', 'SI prey mass', 'Geographic location', 'Latitude', 'Latitude Degree',  

'Latitude label', 'Longitude', 'Longitude Degree', 'Longitude label', 'Longitude Minute',  

'Depth', 'Mean annual temp', 'SD annual temp', 'Mean PP', 'SD PP', 'Specific habitat']]
```

2.2.1 Predator Dataframe

```

In [9]: df_predator['Predator length unit'].value_counts()

Out[9]: mm      22069
        cm      12423
        µm       439
        Name: Predator length unit, dtype: int64

In [10]: for i in range(df_predator.shape[0]):
        if df_predator['Predator length unit'][i] == 'mm':
            df_predator['Predator length'][i] = df_predator['Predator length']
        if df_predator['Predator length unit'][i] == 'µm':
            df_predator['Predator length'][i] = df_predator['Predator length']
        else:
            continue

/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/5563
51888.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df_predator['Predator length'][i] = df_predator['Predator length'][i]*0.1
/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/5563
51888.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df_predator['Predator length'][i] = df_predator['Predator length'][i]*0.0001

In [11]: df_predator['Predator mass unit'].value_counts()

Out[11]: g      34931
        Name: Predator mass unit, dtype: int64

In [12]: df_predator.drop(['Predator length unit'], axis = 1, inplace = True)

/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/2564
476634.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df_predator.drop(['Predator length unit'], axis = 1, inplace = True)

In [13]: df_predator['Predator common name'].value_counts()

Out[13]: Albacore      3581
        Spurdog / spiny dogfish      3287
        Atlantic cod      2518
        Yellowfin tuna      2113
        Atlantic bluefin tuna      1909
        ...
        Tadpole sculpin      2
        4-spot megrim      1
        tope      1
        Imperial scaldfish      1
        Red Irish lord      1
        Name: Predator common name, Length: 87, dtype: int64

In [14]: major_predator = df_predator['Predator common name'].value_counts()/df
major_predator_names = major_predator[major_predator == True].index
major_predator_names

Out[14]: Index(['Albacore', 'Spurdog / spiny dogfish', 'Atlantic cod', 'Yellow
fin tuna',
        'Atlantic bluefin tuna', 'Bigeye tuna'],
        dtype='object')

In [15]: df_names = []
        for name in major_predator_names:
            df_names.append(df_predator[df_predator['Predator common name'] ==
major_predator_new = pd.concat(df_names, ignore_index= True)

In [16]: predator_y = df_predator_new['Predator common name']
        predator_x = df_predator_new[['Predator taxon', 'Predator lifestage', '
Diet coverage', 'Geographic location', 'Dep

```

2.2.2 Prey Dataframe

```
In [17]: df_prey['Prey length unit'].value_counts()
```

```
Out[17]: mm      22471  
cm       11489  
µm        969  
Name: Prey length unit, dtype: int64
```

```
In [18]: for i in range(df_prey.shape[0]):  
    if df_prey['Prey length unit'][i] == 'mm':  
        df_prey['Prey length'][i] = df_prey['Prey length'][i]*0.1  
    if df_prey['Prey length unit'][i] == 'µm':  
        df_prey['Prey length'][i] = df_prey['Prey length'][i]*0.0001  
    else:  
        continue
```

/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/946970684.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
(https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_prey['Prey length'][i] = df_prey['Prey length'][i]*0.1  
/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/946970684.py:5: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
(https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_prey['Prey length'][i] = df_prey['Prey length'][i]*0.0001
```

```
In [19]: df_prey['Prey mass unit'].value_counts()
```

```
Out[19]: g      34728  
mg       203  
Name: Prey mass unit, dtype: int64
```

```
In [20]: for i in range(df_prey.shape[0]):  
    if df_prey['Prey mass unit'][i] == 'mg':  
        df_prey['Prey mass'][i] = df_prey['Prey mass'][i]*0.001  
    else:  
        continue
```

/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/1295962586.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
(https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_prey['Prey mass'][i] = df_prey['Prey mass'][i]*0.001
```

```
In [21]: df_prey.drop(['Prey length unit', 'Prey mass unit'], axis = 1, inplace
```

/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/3688790300.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
(https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

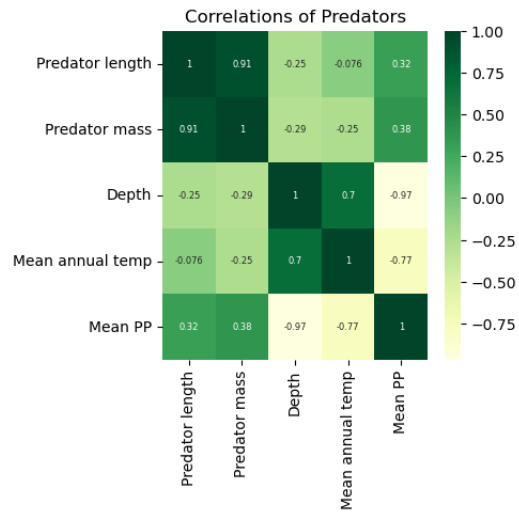
```
df_prey.drop(['Prey length unit', 'Prey mass unit'], axis = 1, inplace = True)
```

```
In [22]: prey_y = df_prey['Prey taxon']  
prey_x = df_prey[['Prey length', 'Prey mass', 'Geographic location', 'De
```


3. Data Visualization

3.1 Correlations of Independent Variables

```
In [23]: # Predators
corr_matrix = predator_x.corr()
fig = plt.figure(figsize = (4,4))
plt.title('Correlations of Predators')
sns.heatmap(corr_matrix, annot=True, annot_kws={"size":6}, center=0, c
plt.show()
```



```
In [24]: # Preys
corr_matrix = prey_x.corr()
fig = plt.figure(figsize = (4,4))
plt.title('Correlations of Preys')
sns.heatmap(corr_matrix, annot=True, annot_kws={"size":6}, center=0, c
plt.show()
```




```
In [25]: df['Prey taxon'].unique()
```


```
Out[25]: array(['mixed', 'ectotherm vertebrate', 'ectotherm invertebrate',
'microzooplankton', 'egg'], dtype=object)
```

3.2 Distribution of Dependent Variable

```
In [26]: fig, ax=plt.subplots(figsize=(5,5))
sns.histplot(predator_y, ax=ax)
plt.xticks(rotation = 90)
plt.show()
```



```
In [27]: fig, ax=plt.subplots(figsize=(5,5))
sns.histplot(pre_y, ax=ax)
plt.xticks(rotation = 90)
plt.show()
```



3.3 Numerial Features

```
In [68]: #predator and prey
df_v = pd.concat([predator_x,prey_x],axis=1)

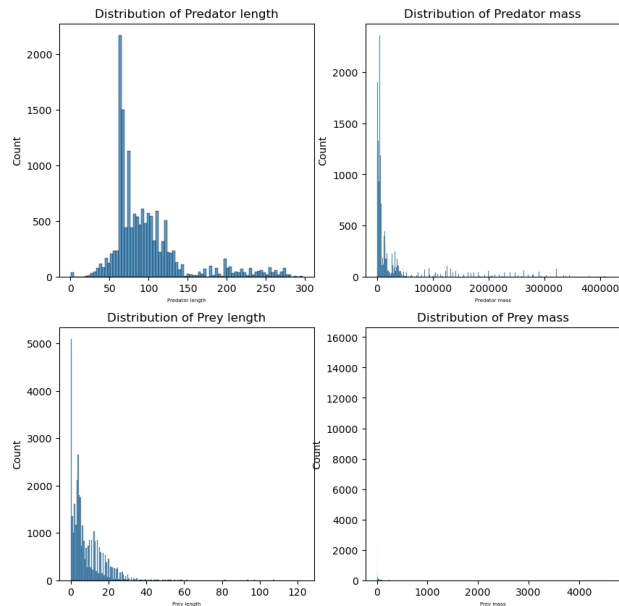
numeric_cols = ['Predator length', 'Predator mass', 'Prey length', 'Prey mass']
print(len(numeric_cols))

fig, ax=plt.subplots(nrows=2, ncols=2, figsize=(10,10))

for var, subplot in zip(numeric_cols, ax.flatten()):
    b=sns.histplot(x=var, data=df_v, ax=subplot)
    b.set_xlabel(str(var), fontsize = 5)
    b.set_title(f'Distribution of {var}')

plt.show()
```

4



```
In [76]: # geographical
df_v = pd.concat([predator_x,prey_x],axis=1)


numeric_cols = ['Depth', 'Mean annual temp', 'Mean PP']
print(len(numeric_cols))

fig, ax=plt.subplots(nrows=1, ncols=3, figsize=(17,6))

for var, subplot in zip(numeric_cols, ax.flatten()):
    b=sns.histplot(x=var, data=df, ax=subplot)
    b.set_xlabel(str(var), fontsize = 10)
    b.set_title(f'Distribution of {var}')

plt.show()
```

3




3.4 Categorical Features


```
In [77]: # Prey Common Names
def myformat(value):
    return f'{value:.1f}%'

threshold = 1000
y = df['Prey common name'].value_counts()
mylabels = y.index
print(mylabels[y >= threshold])
small_values = y[y < threshold]
other_value = sum(small_values)
y = y[y >= threshold]
mylabels = y.index
y = np.append(y, other_value)
mylabels = np.append(mylabels, "Other")

fig1, ax1 = plt.subplots(figsize = (5,5))
ax1.pie(y, labels=mylabels, textprops={'rotation': 0}, autopct=myformat)
ax1.set_title("Prey Common Name distribution")
plt.show()
```


Index(['fish', 'unidentified invertebrate', 'Pearlside', 'copepod', 'mysid', 'Sand lance', 'crustacean'],
dtype='object')



```
In [78]: # Prey taxon
def myformat(value):
    return f'{value:.1f}%'

threshold = 1000
y = df['Prey taxon'].value_counts()
mylabels = y.index

fig1, ax1 = plt.subplots(figsize = (5,5))
ax1.pie(y, labels=mylabels, textprops={'rotation': 0}, autopct=myformat)
ax1.set_title("Prey taxon distribution")
plt.show()
```



```
In [79]: # predator common names
def myformat(value):
    return f'{value:.1f}%'

threshold = 1000
y = df['Predator common name'].value_counts()
mylabels = y.index
print(mylabels[y >= threshold])
small_values = y[y < threshold]
other_value = sum(small_values)
y = y[y >= threshold]
mylabels = y.index
y = np.append(y, other_value)
```

```

mylabels = np.append(mylabels, "Other")

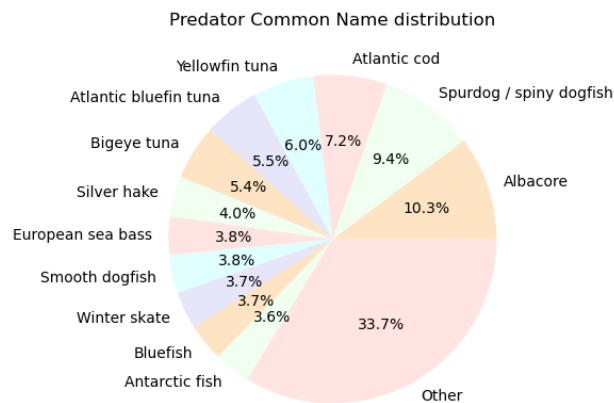
fig1, ax1 = plt.subplots(figsize = (5,5))
ax1.pie(y, labels=mylabels, textprops={'rotation': 0}, autopct=myforma
ax1.set_title("Predator Common Name distribution")
plt.show()

```

```

Index(['Albacore', 'Spurdog / spiny dogfish', 'Atlantic cod', 'Yellow
fin tuna',
      'Atlantic bluefin tuna', 'Bigeye tuna', 'Silver hake',
      'European sea bass', 'Smooth dogfish', 'Winter skate', 'Bluefi
sh',
      'Antarctic fish'],
      dtype='object')

```



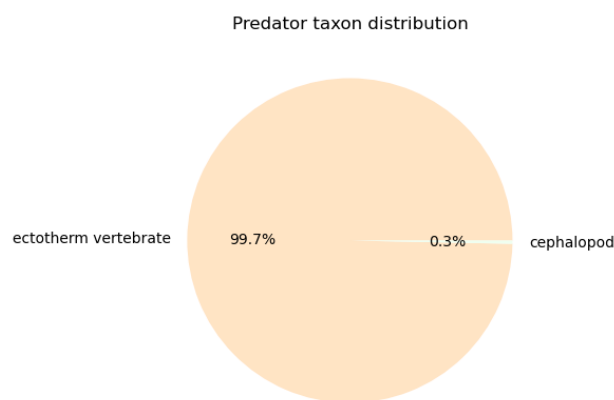
```

In [80]: # Predator taxon
def myformat(value):
    return f'{value:.1f}%'

threshold = 1000
y = df['Predator taxon'].value_counts()
mylabels = y.index

fig1, ax1 = plt.subplots(figsize = (5,5))
ax1.pie(y, labels=mylabels, textprops={'rotation': 0}, autopct=myforma
ax1.set_title("Predator taxon distribution")
plt.show()

```



4. Classification Model

4.1 Classify Predator common names

```

In [81]: # Categorical Feature Encoding
ordinalencoder = OrdinalEncoder()
predator_x['Predator taxon'] = ordinalencoder.fit_transform(predator_x
predator_x['Predator lifestage'] = ordinalencoder.fit_transform(predat
predator_x['Type of feeding interaction'] = ordinalencoder.fit_transfo
predator_x['Diet coverage'] = ordinalencoder.fit_transform(predator_x[
predator_x['Cephalopod'] = ordinalencoder.fit_transform(predator_x[

```

```
predator_x['Geographic location'] = ordinalencoder.fit_transform(predator_x[['Geographic location']])
predator_x['Specific habitat'] = ordinalencoder.fit_transform(predator_x[['Specific habitat']])
```

```
/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/1166
39362.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
predator_x['Predator taxon'] = ordinalencoder.fit_transform(predator_x[['Predator taxon']]).reshape(1,-1).tolist()[0]
/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/1166
39362.py:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
predator_x['Predator lifestage'] = ordinalencoder.fit_transform(predator_x[['Predator lifestage']]).reshape(1,-1).tolist()[0]
/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/1166
39362.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
predator_x['Type of feeding interaction'] = ordinalencoder.fit_transform(predator_x[['Type of feeding interaction']]).reshape(1,-1).tolist()[0]
/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/1166
39362.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
predator_x['Diet coverage'] = ordinalencoder.fit_transform(predator_x[['Diet coverage']]).reshape(1,-1).tolist()[0]
/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/1166
39362.py:7: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
predator_x['Geographic location'] = ordinalencoder.fit_transform(predator_x[['Geographic location']]).reshape(1,-1).tolist()[0]
/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/1166
39362.py:8: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
predator_x['Specific habitat'] = ordinalencoder.fit_transform(predator_x[['Specific habitat']]).reshape(1,-1).tolist()[0]
```

```
In [82]: # Train Test Split
X_train_predator, X_test_predator, y_train_predator, y_test_predator =

# Scaling
scaler = StandardScaler()
X_train_predator = scaler.fit_transform(X_train_predator)
X_test_predator = scaler.transform(X_test_predator)
```

4.1.1 SVC

```
In [83]: svc = SVC(C=1,gamma='scale')
svc.fit(X_train_predator,y_train_predator)
print(f"The train score is:",svc.score(X_train_predator,y_train_predator))
print(f"The test score is:",svc.score(X_test_predator,y_test_predator))

The train score is: 0.893172165958837
The test score is: 0.8919007184846506
```

```
In [84]: # Grid Search
params = {'C': [0.1, 1, 10], 'gamma': [1, 0.1, 0.01]}
svc_gscv = GridSearchCV(estimator = SVC(random_state=123), param_grid=
svc_gscv.fit(X_train_predator, y_train_predator)
print(f'svc best hyperparams : {svc_gscv.best_params_}')
print(f'svc best mean cv accuracy : {svc_gscv.best_score_:.2f}')

svc best hyperparams : {'C': 10, 'gamma': 1}
svc best mean cv accuracy : 0.98
```


```
In [85]: svc_best = SVC(C=10,gamma=1,kernel='linear')
svc_best.fit(X_train_predator,y_train_predator)
```

```
Out[85]: SVC(C=10, gamma=1, kernel='linear')
In a Jupyter environment, please rerun this cell to show the HTML representation or
trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page
with nbviewer.org.
```

```
In [86]: importances = svc_best.coef_[0]
features_dict = dict(zip(predator_x.columns, abs(importances)))
important_features = sorted(features_dict.items(), key=lambda x: -x[1])
important_features
```

```
Out[86]: [('Predator lifestage', 0.2507300387024649),
('Mean PP', 0.2360341213344168),
('Depth', 0.23191475135497236),
('Predator length', 0.13364966017144195),
('Specific habitat', 0.11552890874742702),
('Geographic location', 0.08890184213673946),
('Mean annual temp', 0.07868244720355236),
('Predator mass', 0.06816983848531796),
('Type of feeding interaction', 0.00043798373543962876),
('Predator taxon', 0.0),
('Diet coverage', 0.0)]
```

```
In [87]: sns.barplot(x=predator_x.columns, y=abs(importances))
plt.xticks(rotation=90)
plt.title("Feature Importance of SVC")
plt.xlabel("Feature names")
plt.ylabel("Feature importance")
plt.show()
```



4.1.2 Decision Tree

```
In [88]: dt = DecisionTreeClassifier(criterion='gini', max_depth = 5)
dt.fit(X_train_predator,y_train_predator)
print(f"The train score is:",dt.score(X_train_predator,y_train_predator))
print(f"The test score is:",dt.score(X_test_predator,y_test_predator))
```

The train score is: 0.8951323097027115
The test score is: 0.8909209666884389

```
In [89]: # Grid Search
params = {'criterion': ['gini', 'entropy'], 'max_depth' : [2,3,4,5,6]}
dt_gscv = GridSearchCV(estimator = DecisionTreeClassifier(random_state=42),
                        param_grid=params)
dt_gscv.fit(X_train_predator, y_train_predator)
print(f'dt best hyperparams      : {dt_gscv.best_params_}')
print(f'dt best mean cv accuracy : {dt_gscv.best_score_:.2f}')
```

dt best hyperparams : {'criterion': 'entropy', 'max_depth': 6}
dt best mean cv accuracy : 0.95


```
In [90]: dt_best = DecisionTreeClassifier(criterion='entropy',max_depth=6)
dt_best.fit(X_train_predator, y_train_predator)
```

Out[90]: DecisionTreeClassifier(criterion='entropy', max_depth=6)
In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [91]: importances = dt_best.feature_importances_
features_dict = dict(zip(predator_x.columns, abs(importances)))
important_features = sorted(features_dict.items(), key=lambda x: -x[1])
important_features
```

Out[91]: [('Mean annual temp', 0.4216168393869405),
('Mean PP', 0.3804080518866665),
('Predator mass', 0.10237907590256595),
('Predator length', 0.09559603282382692),
('Predator taxon', 0.0),
('Predator lifestage', 0.0),
('Type of feeding interaction', 0.0),
('Diet coverage', 0.0),
('Geographic location', 0.0),
('Depth', 0.0),
('Specific habitat', 0.0)]

```
In [92]: sns.barplot(x=predator_x.columns, y=abs(importances))
plt.xticks(rotation=90)
plt.title("Feature Importance of Decision Tree")
plt.xlabel("Feature names")
plt.ylabel("Feature importance")
plt.show()
```



4.1.3 Random Forest

```
In [93]: # normal random forest
rfc = RandomForestClassifier(n_estimators=50, max_depth=5)
rfc.fit(X_train_predator, y_train_predator)
print(f"The train score is:", rfc.score(X_train_predator, y_train_predator))
print(f"The test score is:", rfc.score(X_test_predator, y_test_predator))
```

The train score is: 0.9027278667102254
The test score is: 0.9036577400391901

```
In [94]: # cross validation
rfc_cv_scores = cross_val_score(rfc, X_train_predator, y_train_predator, cv=5)
rfc_cv_scores
```

Out[94]: array([0.89383422, 0.89914251, 0.89750919, 0.89873418, 0.90155229])

```
In [95]: depths = [2,4,6,8,10]
train_scores, test_scores = validation_curve(RandomForestClassifier(n_estimators=50,
                                                                    X_train_predator, y_train_predator,
                                                                    param_name='max_depth', param_values=depths))
mean_train_scores = np.average(train_scores, axis=1)
mean_test_scores = np.average(test_scores, axis=1)
pd.DataFrame([mean_train_scores.round(2), mean_test_scores.round(2)],
              columns=pd.Series(depths, name='max_depth'),
              index=['mean_train_scores', 'mean_test_scores'])
```

Out[95]:

	max_depth	2	4	6	8	10
mean_train_scores	0.72	0.87	0.91	0.97	1.00	
mean_test_scores	0.72	0.87	0.91	0.96	0.99	

```
In [96]: # Grid Search
params = {'n_estimators': [10, 50, 100, 150], 'max_depth': [2, 3, 4, 5, 6]}
rfc_gscv = GridSearchCV(estimator=RandomForestClassifier(random_state=42),
                        param_grid=params)
rfc_gscv.fit(X_train_predator, y_train_predator)
print(f'rfc best hyperparams : {rfc_gscv.best_params_}')
print(f'rfc best mean cv accuracy : {rfc_gscv.best_score_:.2f}')
```

rfc best hyperparams : {'max_depth': 6, 'n_estimators': 10}
rfc best mean cv accuracy : 0.92


```
In [110]: rfc_best = RandomForestClassifier(max_depth=6, n_estimators=10)
rfc_best.fit(X_train_predator, y_train_predator)
```

Out[110]: RandomForestClassifier(max_depth=6, n_estimators=10)
In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [111]: importances = rfc_best.feature_importances_
features_dict = dict(zip(predator_x.columns, abs(importances)))
important_features = sorted(features_dict.items(), key=lambda x: -x[1])
important_features
```

Out[111]: [('Specific habitat', 0.256560492030104), ('Geographic location', 0.15289079064402605), ('Predator mass', 0.13723529462744416), ('Predator lifestage', 0.09976287851650764), ('Depth', 0.09067220494493555), ('Mean annual temp', 0.08455929899678413), ('Mean PP', 0.08292953223382424), ('Predator length', 0.08200977517086039), ('Type of feeding interaction', 0.013111345938932368), ('Diet coverage', 0.0002683868965814759), ('Predator taxon', 0.0)]

```
In [112]: sns.barplot(x=predator_x.columns, y=abs(importances))
plt.xticks(rotation=90)
plt.title("Feature Importance of Random Forest")
plt.xlabel("Feature names")
plt.ylabel("Feature importance")
plt.show()
```



4.1.4 Adaboost

```
In [113]: ada = AdaBoostClassifier(n_estimators=50, learning_rate=1)
ada.fit(X_train_predator, y_train_predator)
print(f"The train score is:", ada.score(X_train_predator, y_train_predator))
print(f"The test score is:", ada.score(X_test_predator, y_test_predator))
```

The train score is: 0.4780300555374061
The test score is: 0.48171129980404964

```
In [114]: # Grid Search
params = {'n_estimators':[10,50,100,150], 'learning_rate':[0.01,0.1,0.5]}
ada_gscv = GridSearchCV(estimator=AdaBoostClassifier(random_state=123),
                        param_grid=params)
ada_gscv.fit(X_train_predator, y_train_predator)
print(f'ada best hyperparams : {ada_gscv.best_params_}')
print(f'ada best mean cv accuracy : {ada_gscv.best_score_:.2f}')
```

ada best hyperparams : {'learning_rate': 0.5, 'n_estimators': 50}
ada best mean cv accuracy : 0.83

```
In [115]: ada_best = AdaBoostClassifier(learning_rate=0.5, n_estimators=50)
ada_best.fit(X_train_predator, y_train_predator)
```

Out[115]: AdaBoostClassifier(learning_rate=0.5)


In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [116]: importances = ada_best.feature_importances_
features_dict = dict(zip(predator_x.columns, abs(importances)))
important_features = sorted(features_dict.items(), key=lambda x: -x[1])
important_features
```

Out[116]: [('Specific habitat', 0.26),
('Geographic location', 0.18),
('Mean PP', 0.18),
('Predator length', 0.14),
('Predator mass', 0.08),
('Depth', 0.08),
('Diet coverage', 0.04),
('Predator lifestage', 0.02),
('Mean annual temp', 0.02),
('Predator taxon', 0.0),
('Type of feeding interaction', 0.0)]

```
In [117]: sns.barplot(x=predator_x.columns, y=abs(importances))
plt.xticks(rotation=90)
plt.title("Feature Importance of AdaBoost")
plt.xlabel("Feature names")
plt.ylabel("Feature importance")
plt.show()
```



4.2 Classify Prey taxon

```
In [118]: ordinalencoder = OrdinalEncoder()
prey_x['Geographic location'] = ordinalencoder.fit_transform(pre_x[['Geographic location']])
prey_x['Specific habitat'] = ordinalencoder.fit_transform(pre_x[['Specific habitat']])

/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/3909
440097.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
prey_x['Geographic location'] = ordinalencoder.fit_transform(pre_x[['Geographic location']])
prey_x['Specific habitat'] = ordinalencoder.fit_transform(pre_x[['Specific habitat']])
/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/3909
440097.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
prey_x['Specific habitat'] = ordinalencoder.fit_transform(pre_x[['Specific habitat']])
prey_x['Specific habitat'] = ordinalencoder.fit_transform(pre_x[['Specific habitat']])
```

```
In [119]: from sklearn.model_selection import train_test_split

# Train Test Split
X_train_pre, X_test_pre, y_train_pre, y_test_pre = train_test_split(X, y, random_state=42)

# Scaling
scaler = StandardScaler()
X_train_pre = scaler.fit_transform(X_train_pre)
X_test_pre = scaler.transform(X_test_pre)
```

4.2.1 SVC

```
In [120]: svc = SVC(C=1, gamma='scale')
svc.fit(X_train_pre, y_train_pre)
print(f"The train score is: {svc.score(X_train_pre, y_train_pre)}")
print(f"The test score is: {svc.score(X_test_pre, y_test_pre)}")

The train score is: 0.8626896650443745
The test score is: 0.8605982539001001
```



```
In [121]: # Grid Search
params = {'C': [0.1, 1, 10], 'gamma': [1, 0.1, 0.01]}
svc_gscv = GridSearchCV(estimator = SVC(random_state=123), param_grid=
svc_gscv.fit(X_train_pre, y_train_pre)
print(f'svc best hyperparams      : {svc_gscv.best_params_}')
print(f'svc best mean cv accuracy : {svc_gscv.best_score_:.2f}')

/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-pac
kages/joblib/externals/loky/process_executor.py:702: UserWarning: A w
orker stopped while some jobs were given to the executor. This can be
caused by a too short worker timeout or by a memory leak.
  warnings.warn(

svc best hyperparams      : {'C': 10, 'gamma': 1}
svc best mean cv accuracy : 0.93
```

```
In [122]: svc_best = SVC(C=10,gamma=1,kernel='linear')
svc_best.fit(X_train_pre,y_train_pre)
```

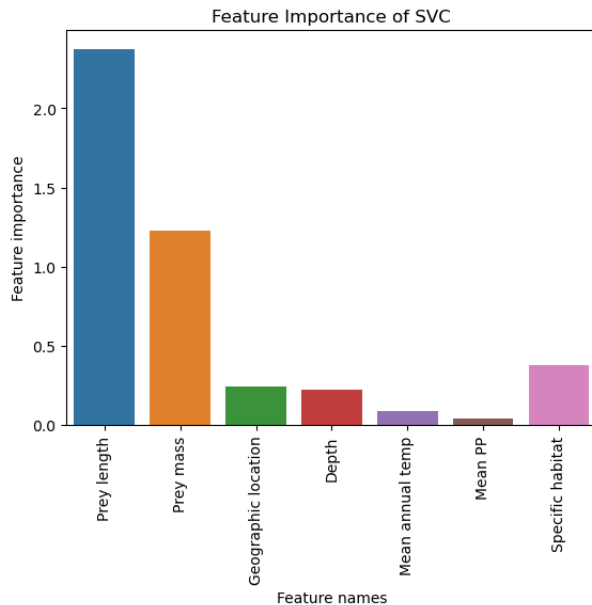
```
Out[122]: SVC(C=10, gamma=1, kernel='linear')

In a Jupyter environment, please rerun this cell to show the HTML representation or
trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page
with nbviewer.org.
```

```
In [123]: importances = svc_best.coef_[0]
features_dict = dict(zip(pre, x.columns, abs(importances)))
important_features = sorted(features_dict.items(), key=lambda x: -x[1])
important_features
```

```
Out[123]: [('Prey length', 2.37413915950674),
('Prey mass', 1.230711912501647),
('Specific habitat', 0.3780779543212702),
('Geographic location', 0.24032763379182143),
('Depth', 0.22540354848024435),
('Mean annual temp', 0.08728371815595892),
('Mean PP', 0.04359612810458202)]
```

```
In [124]: sns.barplot(x=pre, y=abs(importances))
plt.xticks(rotation=90)
plt.title("Feature Importance of SVC")
plt.xlabel("Feature names")
plt.ylabel("Feature importance")
plt.show()
```



4.2.2 Decision Tree

```
In [125]: dt = DecisionTreeClassifier(criterion='gini', max_depth = 5)
dt.fit(X_train_pre,y_train_pre)
print(f"The train score is:",dt.score(X_train_pre,y_train_pre))
print(f"The test score is:",dt.score(X_test_pre,y_test_pre))
```

```
The train score is: 0.8502004008016032
The test score is: 0.8512952626305997
```

```
In [133]: # Grid Search
params = {'criterion': ['entropy'], 'max_depth': [2,3,4,5,6]}
dt_gscv = GridSearchCV(estimator = DecisionTreeClassifier(random_state=0))
dt_gscv.fit(X_train_pre, y_train_pre)
print(f'decision tree best hyperparams      : {dt_gscv.best_params_}')
print(f'decision tree best mean cv accuracy : {dt_gscv.best_score_:.2f}')

decision tree best hyperparams      : {'criterion': 'entropy', 'max_depth': 6}
decision tree best mean cv accuracy : 0.90
```

```
In [134]: dt_best = DecisionTreeClassifier(criterion='entropy',max_depth=6)
dt_best.fit(X_train_pre, y_train_pre)
```

Out[134]: DecisionTreeClassifier(criterion='entropy', max_depth=6)

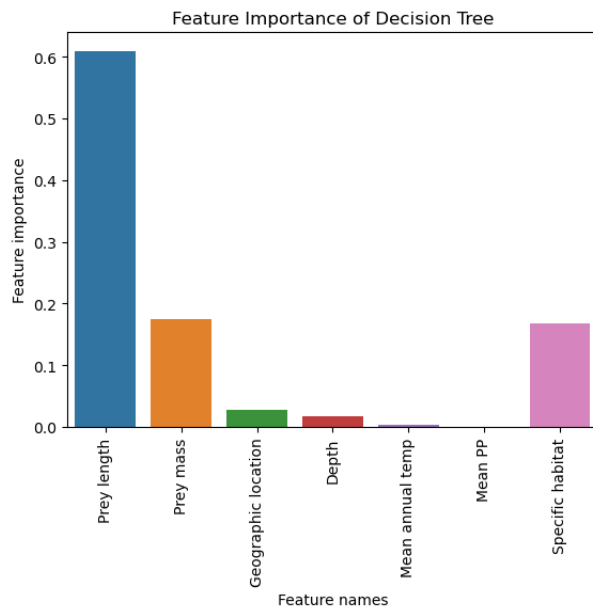
In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [135]: importances = dt_best.feature_importances_
features_dict = dict(zip(pre,abs(importances)))
important_features = sorted(features_dict.items(), key=lambda x: -x[1])
important_features
```

Out[135]: [('Prey length', 0.609117289870658), ('Prey mass', 0.17496708325855898), ('Specific habitat', 0.16818707327772858), ('Geographic location', 0.027699601547847008), ('Depth', 0.017340218522859562), ('Mean annual temp', 0.0026863035722660858), ('Mean PP', 2.4299500816036974e-06)]

```
In [136]: sns.barplot(x=pre, y=abs(importances))
plt.xticks(rotation=90)
plt.title("Feature Importance of Decision Tree")
plt.xlabel("Feature names")
plt.ylabel("Feature importance")
plt.show()
```



4.2.3 Random Forest

```
In [137]: # normal random forest
rfc = RandomForestClassifier(n_estimators=50, max_depth=5)
rfc.fit(X_train_pre, y_train_pre)
print(f"The train score is:", rfc.score(X_train_pre, y_train_pre))
print(f"The test score is:", rfc.score(X_test_pre, y_test_pre))
```

The train score is: 0.84740910392213
The test score is: 0.8494346643766996

```
In [138]: # cross validation
rfc_cv_scores = cross_val_score(rfc, X_train_pre, y_train_pre, cv=5,
rfc_cv_scores
```

```
Out[138]: array([0.86401861, 0.84183217, 0.85274647, 0.83843264, 0.84574087])
```

```
In [139]: depths = [2,4,6,8,10]
train_scores, test_scores = validation_curve(RandomForestClassifier(n_e
X_train_pre, y_train_pre
param_name='max_depth', pa
mean_train_scores = np.average(train_scores, axis=1)
mean_test_scores = np.average(test_scores, axis=1)
pd.DataFrame([mean_train_scores.round(2), mean_test_scores.round(2)],
columns=pd.Series(depths, name='max_depth'),
index=['mean_train_scores', 'mean_test_scores'])
```

```
Out[139]:
```

	max_depth	2	4	6	8	10
mean_train_scores	0.77	0.83	0.87	0.92	0.96	
mean_test_scores	0.77	0.83	0.87	0.92	0.96	

```
In [140]: # Grid Search
params = {'n_estimators':[10,50,100,150], 'max_depth':[2,3,4,5,6]}
rfc_gscv = GridSearchCV(estimator=RandomForestClassifier(random_state=
rfc_gscv.fit(X_train_pre, y_train_pre)
print(f'random forest best hyperparams : {rfc_gscv.best_params_}')
print(f'random forest best mean cv accuracy : {rfc_gscv.best_score_:.2
```

random forest best hyperparams : {'max_depth': 6, 'n_estimators': 10}
random forest best mean cv accuracy : 0.88

```
In [141]: rfc_best = RandomForestClassifier(max_depth=6, n_estimators=10)
rfc_best.fit(X_train_pre, y_train_pre)
```

```
Out[141]: RandomForestClassifier(max_depth=6, n_estimators=10)


In a Jupyter environment, please rerun this cell to show the HTML representation or
trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page
with nbviewer.org.
```

```
In [142]: importances = rfc_best.feature_importances_
features_dict = dict(zip(pre, x.columns, abs(importances)))
important_features = sorted(features_dict.items(), key=lambda x: -x[1])
important_features
```

```
Out[142]: [('Prey length', 0.42992401439283284),
('Prey mass', 0.19545714292515956),
('Specific habitat', 0.1390129260345091),
('Mean PP', 0.07958215448178453),
('Mean annual temp', 0.06859962154022503),
('Depth', 0.05328283389685144),
('Geographic location', 0.034141306728637494)]
```

```
In [143]: sns.barplot(x=prey_x.columns, y=abs(importances))
plt.xticks(rotation=90)
plt.title("Feature Importance of Random Forest")
plt.xlabel("Feature names")
plt.ylabel("Feature importance")
plt.show()
```



4.2.4 Adaboost

```
In [144]: ada = AdaBoostClassifier(n_estimators=50, learning_rate=1)
ada.fit(X_train_pre, y_train_pre)
print(f"The train score is:", ada.score(X_train_pre, y_train_pre))
print(f"The test score is:", ada.score(X_test_pre, y_test_pre))
```

The train score is: 0.802175780131692
The test score is: 0.8043509374552741

```
In [145]: # Grid Search
params = {'n_estimators':[10,50,100,150], 'learning_rate':[0.01,0.1,0.5]}
ada_gscv = GridSearchCV(estimator=AdaBoostClassifier(random_state=123))
ada_gscv.fit(X_train_pre, y_train_pre)
print(f'ada best hyperparams : {ada_gscv.best_params_}')
print(f'ada best mean cv accuracy : {ada_gscv.best_score_:.2f}')
```

ada best hyperparams : {'learning_rate': 1, 'n_estimators': 10}
ada best mean cv accuracy : 0.80

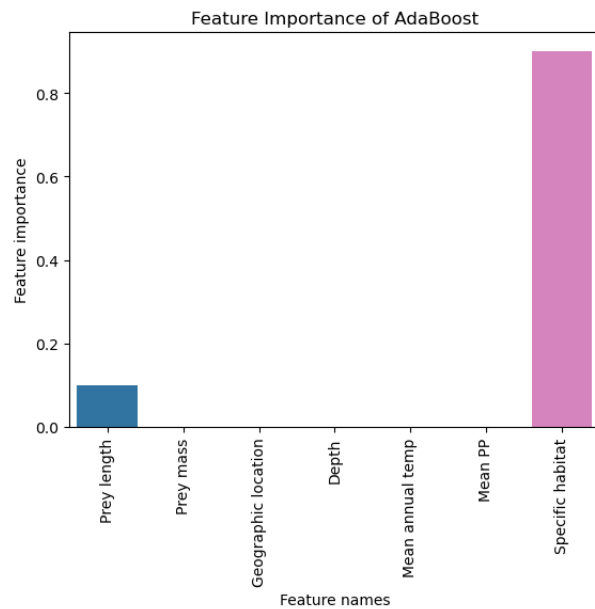
```
In [147]: ada_best = AdaBoostClassifier(learning_rate=1, n_estimators=10)
ada_best.fit(X_train_pre, y_train_pre)
```

Out[147]: AdaBoostClassifier(learning_rate=1, n_estimators=10)
In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [148]: importances = ada_best.feature_importances_
features_dict = dict(zip(pre, abs(importances)))
important_features = sorted(features_dict.items(), key=lambda x: -x[1])
important_features
```

Out[148]: [('Specific habitat', 0.9),
('Prey length', 0.1),
('Prey mass', 0.0),
('Geographic location', 0.0),
('Depth', 0.0),
('Mean annual temp', 0.0),
('Mean PP', 0.0)]

```
In [149]: sns.barplot(x=prey_x.columns, y=abs(importances))
plt.xticks(rotation=90)
plt.title("Feature Importance of AdaBoost")
plt.xlabel("Feature names")
plt.ylabel("Feature importance")
plt.show()
```



5 PCA

5.1 PCA for predator

```
In [150]: pca = PCA(n_components=3, random_state=123)
X_train_pca_predator = pca.fit_transform(X_train_predator)
X_test_pca_predator = pca.transform(X_test_predator)
pca.explained_variance_ratio_
```


```
Out[150]: array([0.3393984 , 0.25719884, 0.19599334])
```

```
In [151]: from mpl_toolkits.mplot3d import Axes3D

fig = plt.figure()
ax = Axes3D(fig)
ax.scatter(X_train_pca_predator[:,0],X_train_pca_predator[:,1],X_train_pca_predator[:,2])
plt.title('PCA for Predator')

/var/folders/8r/f8xs7ssj5375m_ynty13gnw0000gn/T/ipykernel_46169/12821549.py:4: MatplotlibDeprecationWarning: Axes3D(fig) adding itself to the figure is deprecated since 3.4. Pass the keyword argument auto_add_to_figure=False and use fig.add_axes(ax) to suppress this warning.
The default value of auto_add_to_figure will change to False in mpl3.5 and True values will no longer work in 3.6. This is consistent with other Axes classes.
ax = Axes3D(fig)
```

Out[151]: Text(0.5, 0.92, 'PCA for Predator')




```
In [152]: fig, ax=plt.subplots(nrows=1, ncols=3, figsize=(15,5))

sns.scatterplot(x = X_train_pca_predator[:,0], y = X_train_pca_predator[:,2], ax=ax[0], title="Dimensions 0&1 of the PCA Transformation")

sns.scatterplot(x = X_train_pca_predator[:,0], y = X_train_pca_predator[:,1], ax=ax[1], title="Dimensions 0&2 of the PCA Transformation")

sns.scatterplot(x = X_train_pca_predator[:,1], y = X_train_pca_predator[:,2], ax=ax[2], title="Dimensions 1&2 of the PCA Transformation")
```

Out[152]: Text(0.5, 1.0, 'Dimensions 1&2 of the PCA Transformation')



5.2 PCA for prey

```
In [153]: pca = PCA(n_components=3, random_state=123)

X_train_pca_pre = pca.fit_transform(X_train_pre)
X_test_pca_pre = pca.transform(X_test_pre)

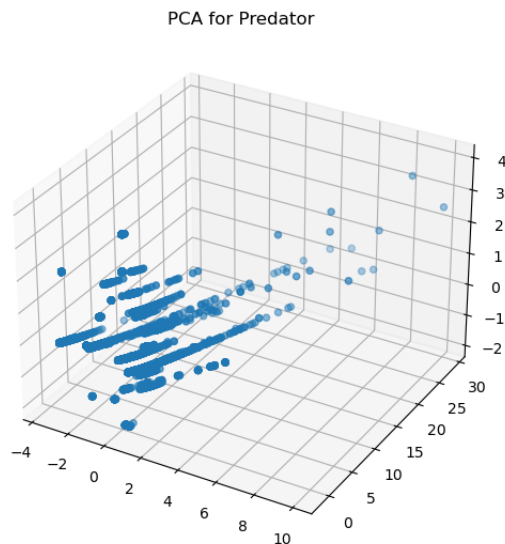
pca.explained_variance_ratio_
```

Out[153]: array([0.34895859, 0.22353736, 0.15422669])

```
In [154]: fig = plt.figure()
ax = Axes3D(fig)
ax.scatter(X_train_pca_pre[:,0],X_train_pca_pre[:,1],X_train_pca_pre[:,2])
plt.title('PCA for Predator')

/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/4104782933.py:2: MatplotlibDeprecationWarning: Axes3D(fig) adding itself to the figure is deprecated since 3.4. Pass the keyword argument auto_add_to_figure=False and use fig.add_axes(ax) to suppress this warning. The default value of auto_add_to_figure will change to False in matplotlib 3.5 and True values will no longer work in 3.6. This is consistent with other Axes classes.
ax = Axes3D(fig)
```

Out[154]: Text(0.5, 0.92, 'PCA for Predator')




```
In [155]: fig, ax=plt.subplots(nrows=1, ncols=3, figsize=(15,5))

sns.scatterplot(x = X_train_pca_pre[:,0], y = X_train_pca_pre[:,1],
ax[0].set_title("Dimensions 0&1 of the PCA Transformation")

sns.scatterplot(x = X_train_pca_pre[:,0], y = X_train_pca_pre[:,2],
ax[1].set_title("Dimensions 0&2 of the PCA Transformation")

sns.scatterplot(x = X_train_pca_pre[:,1], y = X_train_pca_pre[:,2],
ax[2].set_title("Dimensions 1&2 of the PCA Transformation"))
```

Out[155]: Text(0.5, 1.0, 'Dimensions 1&2 of the PCA Transformation')




6 Clustering Model: Kmeans

6.1 Kmeans to cluster predator

```
In [156]: inertia = []
for i in range(1, 11):
    km = KMeans(n_clusters=i, random_state=0)
    km.fit(predator_x)
    inertia.append(km.inertia_)

plt.plot(range(1, 11), inertia, marker='o')
```

Out[156]: [



```
In [157]: kmeans_predator = KMeans(n_clusters=2, random_state=123)
kmeans_predator.fit(predator_x)
labels_predator = kmeans_predator.labels_
labels_predator = pd.DataFrame(labels_predator, columns=["kmeans_label"])
labels_predator.value_counts()
```

Out[157]: kmeans_label
0 14082
1 1224
dtype: int64

```
In [158]: X_predator_cluster = pd.concat([predator_x, predator_y, labels_predator], axis=1)
X_predator_cluster.groupby('Predator common name').mean()
```

Out[158]:

	Predator taxon	Predator lifestage	Type of feeding interaction	Predator length	Predator mass	Diet coverage	Geographic location	
Predator common name								
Albacore	0.0	1.000000	0.336219	66.883552	6096.995309	0.000000	9.000000	4
Atlantic bluefin tuna	0.0	0.000000	0.123625	203.665584	156442.529544	0.000000	6.643793	
Atlantic cod	0.0	0.040111	0.754170	66.486644	2691.241660	0.181493	0.209293	
Bigeye tuna	0.0	0.000000	0.804004	107.865121	14839.720179	0.000000	3.000000	4
Spurdog / spiny dogfish	0.0	0.000000	0.421053	82.239428	2547.490880	0.031031	0.020688	
Yellowfin tuna	0.0	0.000000	1.228585	108.681022	26764.666540	0.000000	3.000000	4

```
In [159]: y_train_predator.value_counts()
```


Out[159]: Albacore 2895
Spurdog / spiny dogfish 2609
Atlantic cod 2003
Yellowfin tuna 1712
Atlantic bluefin tuna 1514
Bigeye tuna 1511
Name: Predator common name, dtype: int64

6.2 Kmeans to cluster prey


```
In [160]: inertia = []
for i in range(1, 11):
    km = KMeans(n_clusters=i, random_state=0)
    km.fit(pre_x)
    inertia.append(km.inertia_)

plt.plot(range(1, 11), inertia, marker='o')
```

Out[160]: [



```
In [161]: kmeans_pre = KMeans(n_clusters=2, random_state=123)
kmeans_pre.fit(pre_x)
labels_pre = kmeans_pre.labels_
labels_pre = pd.DataFrame(labels_pre, columns=["kmeans_label"])
labels_pre.value_counts()
```

Out[161]: kmeans_label
1 26012
0 8919
dtype: int64

```
In [162]: X_pre_cluster = pd.concat([pre_x, pre_y, labels_pre], axis=1)
X_pre_cluster.groupby('Prey taxon').mean()
```

Out[162]:

	Prey length	Prey mass	Geographic location	Depth	Mean annual temp	Mean PP	Sp h
Prey taxon							
ectotherm invertebrate	3.560657	13.905011	8.013103	1599.555928	14.518936	726.588607	8.1
ectotherm vertebrate	11.430129	39.217246	7.630517	1531.068544	15.756414	739.622286	6.9
egg	0.030330	0.000002	5.080944	771.573356	1.017032	509.966273	7.8
microzooplankton	0.060302	0.000008	14.585779	2559.448081	18.395485	562.847630	16.5
mixed	11.351194	21.888489	4.559006	14.006211	23.652174	866.000000	0.0

7. Regression Model

```
In [163]: merged_df = pd.merge(df_predator, df_pre, left_index=True, right_index=True)

# print(merged_df.columns)

merged_df['Mass difference'] = merged_df['Predator mass'] - merged_df['Prey mass']
# merged_df['Mass difference'] = merged_df['Predator length'] - merged_df['Prey length']

reg_x = merged_df[['Type of feeding interaction', 'Geographic location', 'Latitude Minute_x', 'Latitude label_x', 'Longitude label_x', 'Depth_x', 'Mean annual temp_x', 'Mean PP_x', 'SD PP_x', 'Specific habitat_x']]

reg_y = merged_df['Mass difference']
```

```
In [164]: reg_y
```

```
Out[164]: 0      1525.6260
          1      1591.7787
          2      1831.7070
          3       79.5090
          4       57.3037
          ...
        34926    510.7000
        34927    508.9000
        34928       6.6261
        34929       6.6261
        34930    369.8300
Name: Mass difference, Length: 34931, dtype: float64
```

```
In [165]: non_numeric_columns = reg_x.select_dtypes(exclude=['int64', 'float64'])
ordinalencoder = OrdinalEncoder()

for col_name in list(non_numeric_columns):
    # print(col_name)
    reg_x[col_name] = reg_x[col_name].astype(str)
    reg_x[col_name] = ordinalencoder.fit_transform(reg_x[[col_name]]).

/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/49
3454921.py:7: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pa
das-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-
copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/index
ing.html#returning-a-view-versus-a-copy)
    reg_x[col_name] = reg_x[col_name].astype(str)
/var/folders/8r/f8xs7ssj5375m_ynty13gnww0000gn/T/ipykernel_46169/49
3454921.py:8: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pa
das-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-
copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/index
ing.html#returning-a-view-versus-a-copy)
    reg_x[col_name] = ordinalencoder.fit_transform(reg_x[[col_name]])
```

```
In [166]: # Train Test Split
X_train_len, X_test_len, y_train_len, y_test_len = train_test_split(reg_y, reg_x,
                                                                      test_size=0.2,
                                                                      random_state=42)

# Scaling
scaler = StandardScaler()
X_train_len = scaler.fit_transform(X_train_len)
X_test_len = scaler.transform(X_test_len)
```

7.1 Linear Regression

```
In [167]: reg = LinearRegression()
reg.fit(X_train_len,y_train_len)
print(f"The train score:", reg.score(X_train_len,y_train_len))
print(f"The test score:", reg.score(X_test_len,y_test_len))
```

The train score: 0.5808598730874159
The test score: 0.5563155436980785

7.2 Ridge Regression

```
In [168]: ridge = Ridge(alpha=1)
ridge.fit(X_train_len,y_train_len)
print(f"The train score:", ridge.score(X_train_len,y_train_len))
print(f"The test score:", ridge.score(X_test_len,y_test_len))
```

The train score: 0.5807786367386404
The test score: 0.5561781472999623

```
In [169]: params = {'alpha':[0.01,0.1,0.5,1]}
ridge_gscv = GridSearchCV(estimator=Ridge(), param_grid=params, cv=3,
                           scoring='neg_mean_squared_error')
ridge_gscv.fit(X_train_len, y_train_len)
print(f'ada best hyperparams      : {ridge_gscv.best_params_}')
print(f'ada best mean cv accuracy : {ridge_gscv.best_score_:.2f}')

ada best hyperparams      : {'alpha': 0.1}
ada best mean cv accuracy : 0.58
```

7.3 Lasso Regression

```
In [170]: lasso = Lasso(alpha=1)
lasso.fit(X_train_len,y_train_len)
print(f"The train score:", lasso.score(X_train_len,y_train_len))
print(f"The test score:", lasso.score(X_test_len,y_test_len))

The train score: 0.5787062448058675
The test score: 0.5531628009202632

/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-packages/sklearn/linear_model/_coordinate_descent.py:648: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 1.241e+13, tolerance: 6.062e+09
  model = cd_fast.enet_coordinate_descent(

In [171]: params = {'alpha':[0.01,0.1,0.5,1]}
lasso_gscv = GridSearchCV(estimator=Lasso(), param_grid=params, cv=3,
lasso_gscv.fit(X_train_len, y_train_len)
print(f'ada best hyperparams      : {lasso_gscv.best_params_}')
print(f'ada best mean cv accuracy : {lasso_gscv.best_score_:.2f}')

/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-packages/sklearn/linear_model/_coordinate_descent.py:648: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 8.647e+12, tolerance: 4.066e+09
  model = cd_fast.enet_coordinate_descent(
/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-packages/sklearn/linear_model/_coordinate_descent.py:648: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 8.664e+12, tolerance: 4.118e+09
  model = cd_fast.enet_coordinate_descent(
/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-packages/sklearn/linear_model/_coordinate_descent.py:648: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 8.055e+12, tolerance: 3.940e+09
  model = cd_fast.enet_coordinate_descent(
/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-packages/sklearn/linear_model/_coordinate_descent.py:648: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 8.764e+12, tolerance: 4.118e+09
  model = cd_fast.enet_coordinate_descent(
/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-packages/sklearn/linear_model/_coordinate_descent.py:648: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 8.075e+12, tolerance: 3.940e+09
  model = cd_fast.enet_coordinate_descent(
/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-packages/sklearn/linear_model/_coordinate_descent.py:648: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 8.787e+12, tolerance: 4.118e+09
  model = cd_fast.enet_coordinate_descent(
/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-packages/sklearn/linear_model/_coordinate_descent.py:648: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 8.625e+12, tolerance: 4.066e+09
  model = cd_fast.enet_coordinate_descent(
/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-packages/sklearn/linear_model/_coordinate_descent.py:648: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 7.968e+12, tolerance: 3.940e+09
  model = cd_fast.enet_coordinate_descent(
/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-packages/sklearn/linear_model/_coordinate_descent.py:648: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 8.530e+12, tolerance: 4.066e+09
  model = cd_fast.enet_coordinate_descent(
/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-packages/sklearn/linear_model/_coordinate_descent.py:648: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 8.540e+12, tolerance: 4.118e+09
  model = cd_fast.enet_coordinate_descent(
/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-packages/sklearn/linear_model/_coordinate_descent.py:648: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 7.860e+12, tolerance: 3.940e+09
  model = cd_fast.enet_coordinate_descent(
/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-packages/sklearn/linear_model/_coordinate_descent.py:648: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 8.412e+12, tolerance: 4.066e+09
  model = cd_fast.enet_coordinate_descent(

ada best hyperparams      : {'alpha': 0.01}
ada best mean cv accuracy : 0.58

/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-packages/sklearn/linear_model/_coordinate_descent.py:648: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 1.276e+13, tolerance: 6.062e+09
  model = cd_fast.enet_coordinate_descent(
```

7.5 Decision Tree Regression

```
The train score: 0.7479117611237333
The test score: 0.7489658353079156
```

```

/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-p
ackages/sklearn/tree/_classes.py:397: FutureWarning: Criterion 'mse
' was deprecated in v1.0 and will be removed in version 1.2. Use 'c
riterion='squared_error'' which is equivalent.
  warnings.warn(
/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-p
ackages/sklearn/tree/_classes.py:397: FutureWarning: Criterion 'mse
' was deprecated in v1.0 and will be removed in version 1.2. Use 'c
riterion='squared_error'' which is equivalent.
  warnings.warn(
/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-p
ackages/sklearn/tree/_classes.py:397: FutureWarning: Criterion 'mse
' was deprecated in v1.0 and will be removed in version 1.2. Use 'c
riterion='squared_error'' which is equivalent.
  warnings.warn(
/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-p
ackages/sklearn/tree/_classes.py:397: FutureWarning: Criterion 'mse
' was deprecated in v1.0 and will be removed in version 1.2. Use 'c
riterion='squared_error'' which is equivalent.

```

	index	0
0	Latitude Minute	0.695717
1	Mean annual temp	0.182461
2	Geographic location	0.079401
3	Latitude	0.022632
4	Type of feeding interaction	0.009791
5	SD annual temp	0.006831

Feature	Importance
Latitude Minute	0.68
Mean annual temp	0.18
Geographic location	0.08
Latitude	0.02
Type of feeding interaction	0.01
SD annual temp	0.01

The train score: 0.5799527547426957
The test score: 0.5931023110899205

```
In [177]: params = {'n_estimators':[10,50,100,150], 'learning_rate':[0.01,0.1,0.5]
ada_gscv = GridSearchCV(estimator=AdaBoostRegressor(random_state=123),
ada_gscv.fit(X_train_len, y_train_len)
print(f'ada best hyperparams : {ada_gscv.best_params_}')
print(f'ada best mean cv accuracy : {ada_gscv.best_score_:.2f}')
```

/Users/jingyi/opt/anaconda3/envs/MyCodingSpace/lib/python3.9/site-packages/joblib/externals/loky/process_executor.py:702: UserWarning: A worker stopped while some jobs were given to the executor. This can be caused by a too short worker timeout or by a memory leak.

```
warnings.warn(

ada best hyperparams : {'learning_rate': 0.01, 'n_estimators': 150}
ada best mean cv accuracy : 0.71
```

```
In [178]: df = pd.DataFrame(adarg.feature_importances_, index=reg_x.columns).sort_index()
df['index'] = df['index'].str.split('_').str[0]
df
```

```
Out[178]:
```


	index	0
0	Geographic location	0.318066
1	Latitude Minute	0.277111
2	Latitude	0.202147
3	Latitude Degree	0.086625
4	Longitude Degree	0.067136
5	SD PP	0.024572

```
In [179]: sns.set_style("whitegrid")
sns.set(rc={'figure.figsize':(11.7,5)})

# Create the bar chart using seaborn's barplot function
ax = sns.barplot(y=0, x='index', data=df)

# Set the title and axis labels
ax.set_title("Feature Importance for AdaBoost")
ax.set_xlabel("Importance")
ax.set_ylabel("Feature")

# Show the plot
plt.show()
```



5.6 Random Forest Regression

```
In [181]: rf = RandomForestRegressor(max_depth=5)
rf.fit(X_train_len,y_train_len)
print(f"The train score:", dtr.score(X_train_len,y_train_len))
print(f"The test score:", dtr.score(X_test_len,y_test_len))

depths = [2,4,6,8,10]
train_scores,test_scores = validation_curve(RandomForestRegressor(n_estimators=100,
                                                                    X_train_len, y_train_len,
                                                                    param_name='max_depth', param_values=depths))

mean_train_scores = np.average(train_scores, axis=1)
mean_test_scores = np.average(test_scores, axis=1)
pd.DataFrame([mean_train_scores.round(2),mean_test_scores.round(2)],
              columns=pd.Series(depths,name='max_depth'),
              index=['mean_train_scores','mean_test_scores'])
```

The train score: 0.7479117611237333
The test score: 0.7489658353079156

```
Out[181]:
```

	max_depth	2	4	6	8	10
mean_train_scores	0.41	0.72	0.75	0.75	0.75	0.75
mean_test_scores	0.40	0.72	0.74	0.74	0.74	0.74

```
In [182]: params = {'n_estimators':[10,50,100,150], 'max_depth':[2,3,4,5,6]}
rfc_gscv = GridSearchCV(estimator=RandomForestRegressor(random_state=1),
                        param_grid=params)
rfc_gscv.fit(X_train_len, y_train_len)
print(f'random forest best hyperparams : {rfc_gscv.best_params_}')
print(f'random forest best mean cv accuracy : {rfc_gscv.best_score_:0.2f}')

random forest best hyperparams : {'max_depth': 6, 'n_estimators': 10}
random forest best mean cv accuracy : 0.74
```

```
In [183]: df = pd.DataFrame(rf.feature_importances_, index=reg_x.columns).sort_values(
df['index'] = df['index'].str.split('_').str[0]
df
```

```
Out[183]:
```

	index	0
0	Latitude Minute	0.570342
1	Mean annual temp	0.101127
2	Geographic location	0.095712
3	SD PP	0.049368
4	SD annual temp	0.041918
5	Longitude	0.037413

```
In [184]: sns.set_style("whitegrid")
sns.set(rc={'figure.figsize':(11.7,5)})

# Create the bar chart using seaborn's barplot function
ax = sns.barplot(y=0, x='index', data=df)

# Set the title and axis labels
ax.set_title("Feature Importance for Random Forest")
ax.set_xlabel("Importance")
ax.set_ylabel("Feature")

# Show the plot
plt.show()
```

